

# WS 2021/22 1048 - Econometrics I

## Case Study 4

Anja Kulagic (h12100543), Jovana Mileusnic (h12100542), Ema Vargova (h11914081)

## Contents

<b>1 Data Description</b>	<b>3</b>
<b>2 Model</b>	<b>4</b>
2.1 Model estimation . . . . .	4
2.2 . . . . .	4
2.3 . . . . .	5
2.4 . . . . .	5
2.5 . . . . .	7
2.6 . . . . .	9
2.7 . . . . .	9
<b>3 Theory</b>	<b>9</b>
3.1 . . . . .	9
3.2 . . . . .	10
3.3 . . . . .	11
3.4 . . . . .	12
<b>References</b>	<b>14</b>

## List of Figures

## List of Tables

1	Multiple OLS Regression Models . . . . .	5
2	Testing the null hypothesis: There is no difference in the average rating between the brands rq and wa, ceteris paribus. . . . .	6
3	Testing the null hypothesis: There is no difference in the average rating between the brands rq and wa, ceteris paribus. . . . .	6
4	Testing the null hypothesis: The brand information is not helpful to determine the rating of mineral water. . . . .	8
5	Testing the null hypothesis: The brand information is not helpful to determine the rating of mineral water. . . . .	8
6	Comparing models in terms of model selection criteria . . . . .	8
7	Proof for 3.1. . . . .	10
8	Simulated Regression Model . . . . .	13
9	The range of explanatory variable . . . . .	13
10	Testing the null hypothesis: The vertex of a quadratic equation equals one. . . . .	14

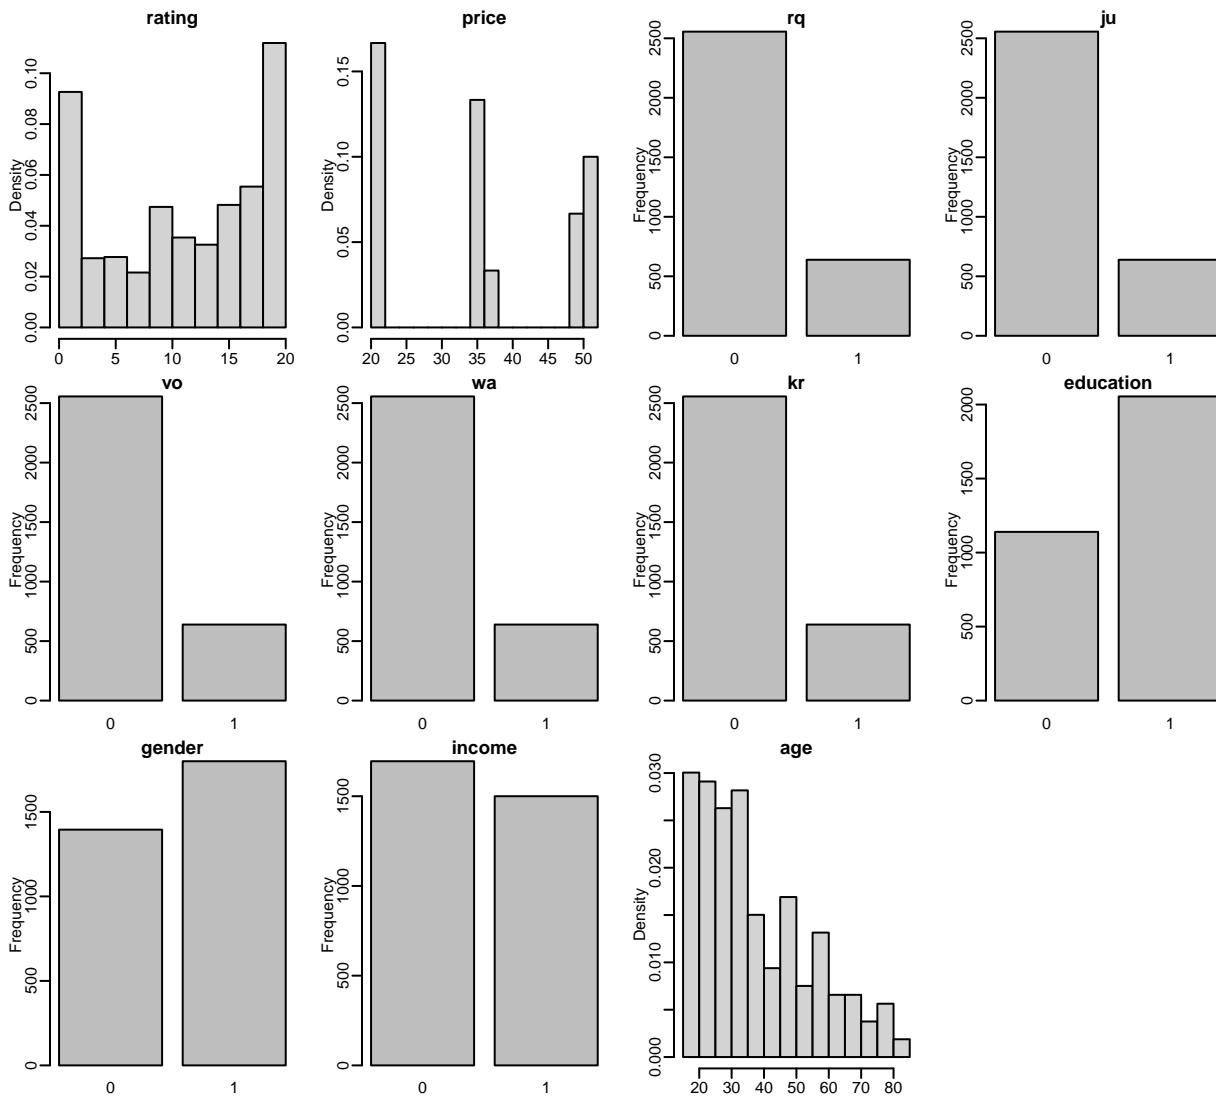
# 1 Data Description

```

data <- read.csv("marketing.csv")
data$rating <- as.numeric(data$rating)
data$age <- as.numeric(data$age )
metric.index <- which(lapply(data, class)=="numeric")
factor.index <- which(lapply(data, class)=="integer")

op <- par(mfrow=c(3, 4), cex=.5, mar = c(1.5, 2.5, 1.5, 1), mgp = c(1.3, .5, 0))
for(i in 1:ncol(data)) {
  if(i %in% metric.index) {
    hist(data[[i]], freq=FALSE, main=names(data)[i], xlab="")
  } else {
    barplot(table(data[[i]]), main=names(data)[i], ylab="Frequency" )}
}

```



The provided dataset contains 11 variables where we can distinguish two groups of dummy variables, along with the variables price and age as one continuous and one discrete variable, respectively.

The first group of dummy variable is related to the brand of mineral water considering five different brands, sorting data into mutually exclusive categories (meaning that each consumer rating should be based on one out of five observed brands of mineral water,  $rq_i + ju_i + vo_i + wa_i + kr_i = 1$ ). This should be taken into consideration if we are building a multiple linear regression model as one of main Gauss-Markov conditions could be violated (*dummy variable trap*). Note that dataset includes 639 observations for each observed brand of mineral water.

However, so-called *zero-one variables* related to education, gender and income are not complementary, as there is not any kind of relationship between them, thus we can face with man consumer with no high school diploma and below average income, meaning that all previously mentioned dummy variables equal zero. If we look briefly at the price and age histograms, we can spot existence of three price categories up to a certain point, whereas we could notice that there is positive asymmetry in age distribution.

## 2 Model

### 2.1 Model estimation

#### 2.1.1

```
model1 <- lm(rating ~ . - rq, data)
```

#### 2.1.2

```
model2 <- lm(rating ~ . - 1, data=data)
```

```
library(stargazer)
stargazer(model1, model2, single.row=TRUE, header=FALSE, model.numbers=FALSE,
           column.labels = c("Model 1", "Model 2"), title="Multiple OLS Regression Models")
```

Table 1 was created using stargazer v.5.2.2 by Hlavac (2018).

In order to avoid so called *dummy variable trap*, solution is to drop one of the categorical variables, or in other words to let one mineral water brand to be a baseline. Using *rq* as a baseline for the brand effect, we consider this brand to the *reference value* while the values of the remaining categories represent the change from this reference. In other words, the brand *rq* stands for the group against which comparison is made. All other variables are considered as self-explanatory variables.

Alternative way to avoid dummy variable trap is to drop the intercept constant from the multiple regression model. In our case, this is illustrated with the model 2, which now does not use a brand as a baseline but includes all brand dummies in the model, while excluding the intercept. However, after comparing OLS estimates of coefficients related to self-explanatory variables from both models, it is clear that this change does not have any influence on their values. To rephrase it, they are exactly the same in both models. Even though including g dummies without an overall intercept is sometimes useful, it has two practical drawbacks which we should be aware of. Firstly, it makes it more cumbersome to test for differences relative to a base group and additionally, we will face with the uncentered R-squared, which will be discussed later.

### 2.2

Model 1: To interpret the coefficients on the dummy variables, we must remember that the base group, or so-called benchmark, is the mineral water brand *rq*. Thus, the estimates on the four dummy variables

Table 1: Multiple OLS Regression Models

	<i>Dependent variable:</i>	
	rating	Model 2
Model 1		
price	-0.303*** (0.008)	-0.303*** (0.008)
rq		24.732*** (0.478)
ju	-3.884*** (0.312)	20.848*** (0.478)
vo	-0.327 (0.312)	24.405*** (0.478)
wa	-3.288*** (0.312)	21.444*** (0.478)
kr	-4.172*** (0.312)	20.560*** (0.478)
education	-0.257 (0.218)	-0.257 (0.218)
gender	-0.107 (0.200)	-0.107 (0.200)
income	-0.641*** (0.205)	-0.641*** (0.205)
age	0.012** (0.006)	0.012** (0.006)
Constant	24.732*** (0.478)	
Observations	3,195	3,195
R <sup>2</sup>	0.348	0.828
Adjusted R <sup>2</sup>	0.346	0.828
Residual Std. Error (df = 3185)	5.584	5.584
F Statistic	188.881*** (df = 9; 3185)	1,537.900*** (df = 10; 3185)

Note:

\*p&lt;0.1; \*\*p&lt;0.05; \*\*\*p&lt;0.01

measure the proportionate difference in consumer rating relative to the brand rq. To be specific, the brand vo is estimated to have 0.327 lower rating than the brand rq, with the same levels of the other variables.

Model 2: Looking at the estimator related to the brand vo, we can conclude that if the dummy variable of the brand vo is equal to one, all other dummy variables related to mineral water brands must be zero. Furthermore, it will have a positive effect of 24.405 on consumer rating, with all other variables being constant.

## 2.3

We can calculate the regression parameter associated with vo in Model 1 (denoted  $\hat{\beta}_3$ ) from the regression parameters of Model 2 (denoted  $\tilde{\beta}_4$ ) in the following way:

$$\begin{aligned}\hat{\beta}_3 &= \tilde{\beta}_2 + \tilde{\beta}_4 \\ \hat{\beta}_3 &= 24.732 - 24.405 \approx 0.327\end{aligned}$$

To be specific, in Model 1 the coefficient on the mineral water brand rq would be represented by the intercept, whereas the estimates on the remaining brands of mineral water would measure the proportionate difference in consumer rating relative to the brand rq. Therefore, the straightforward way of computing the estimate on the brand vo in Model 1 is by subtracting the estimation on vo from the estimation on rq in Model 2.

## 2.4

We are interested in hypothesis that average rating between brands rq and wa are identical, with the same levels of the other variables, at a significance level of  $\alpha = 0.05$ .

For the usual t test to be valid, we must assume that the homoskedasticity assumption holds, which means that the population variance in ratings is the same among the different brands of mineral water, namely, the brand rq and the brand wa.

For Model1:

```
library(car)
```

```
## Warning: package 'car' was built under R version 4.1.2
```

```
kable(linearHypothesis(model1, c("wa=0")),
      caption = "Testing the null hypothesis: There is no difference in the average rating
between the brands rq and wa, ceteris paribus.")
```

Table 2: Testing the null hypothesis: There is no difference in the average rating between the brands rq and wa, ceteris paribus.

Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
3186	102773.54	NA	NA	NA	NA
3185	99319.55	1	3453.991	110.7633	0

Under the null:

$$t = \frac{\hat{\beta}_4}{se(\hat{\beta}_4|X)} | X \sim t_{df}$$

$$t = \frac{-3.288}{0.312} \approx -10.538$$

The estimated difference, -3.288, has a t statistic of -10.538, which is very statistically significant, and p-value of 2e-16 thus we conclude that the estimated difference between the brands rq and wa is very statistically significant and reject the null.

To conclude, when we are doing the testing with dummy variables, the easiest thing to do it to choose one of these groups to be the base group in order to get needed estimates and its standard error directly.

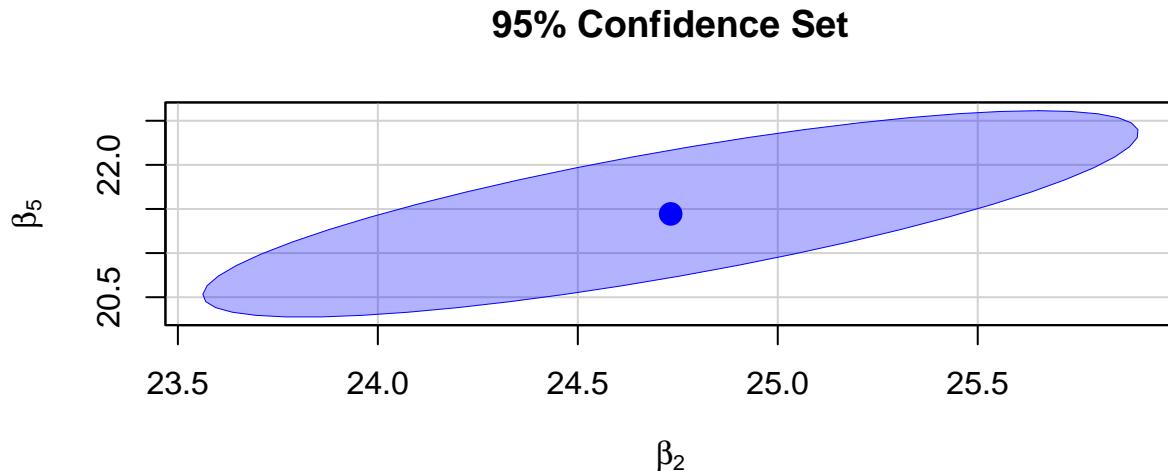
For Model2:

```
kable(linearHypothesis(model2, c("rq=wa")),
      caption = "Testing the null hypothesis: There is no difference in the average rating between the brands rq and wa, ceteris paribus.")
```

Table 3: Testing the null hypothesis: There is no difference in the average rating between the brands rq and wa, ceteris paribus.

Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
3186	102773.54	NA	NA	NA	NA
3185	99319.55	1	3453.991	110.7633	0

```
confidenceEllipse(model2, fill = T, lwd = 0,
                  which.coef = c("rq", "wa"),
                  main = "95% Confidence Set", ylab=bquote(paste(beta)[5]), xlab=bquote(paste(beta)[2]))
```



Running a test on Model 2 would lead us to the same conclusion. On contrary, the F-test would be the preferred method for test in joint hypothesis which requires running two regressions, one regression on the unrestricted model and one regression on the restricted model, which we have imposed the restrictions of the null.

Manually, we can compute the F-statistic as follows:

$$F = \frac{(SSR_R - SSR_{UR})/g}{SSR_{UR}/(n - k - 1)}$$

and the intuition behind it is:

$$F = \frac{\text{Average loss in explanatory power under } H_0}{\text{Average unexplained variation under } H_A}$$

We obtain  $F=110.76$ , which is high value, thus we can conclude that we lose a lot of explanatory power by our restrictions. The 95% confidence set for coefficients on brands rq and wa (namely  $\beta_2$  and  $\beta_5$ ) is an ellipse which contains the pairs of values of  $\beta_2$  and  $\beta_5$  that cannot be rejected using F-statistic at the 5% significance level.

## 2.5

Building a reduced Model 3:

```
model3 <- lm(rating ~ price + education + gender + income + age, data=data)
```

### 2.5.1

```
kable(linearHypothesis(model1, c("ju=0", "vo=0", "wa=0", "kr=0")),
      caption = "Testing the null hypothesis: The brand information
is not helpful to determine the rating of mineral water.")
```

Table 4: Testing the null hypothesis: The brand information is not helpful to determine the rating of mineral water.

Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
3189	109650.39	NA	NA	NA	NA
3185	99319.55	4	10330.84	82.82291	0

```
kable(linearHypothesis(model2, c("rq=0", "ju=0", "vo=0", "wa=0", "kr=0")),
      caption = "Testing the null hypothesis:
The brand information is not helpful to determine the rating of mineral water.")
```

Table 5: Testing the null hypothesis: The brand information is not helpful to determine the rating of mineral water.

Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
3190	192342.18	NA	NA	NA	NA
3185	99319.55	5	93022.63	596.6138	0

In both cases, the obtained results lead us to the same conclusion. For the given p-value, we would reject null hypothesis which states that the brand information is not helpful to determine the rating of mineral water, as it is smaller than any significance level commonly used in practice.

However, we can notice that the obtained F-test differ between Model 1 and Model 2, due to high value of the sum of squared residuals (SSR) of restricted model, which may point out that if we exclude the brand information from our model, we will lose a lot of explanatory power due to our restriction. The loss is mainly noticeable in Model 2 caused by excluding intercept and the brand information which are obviously significant for our Model.

## 2.5.2

Table 6: Comparing models in terms of model selection criteria

	Model 1	Model 2	Model 3
R-squared	0.348	0.828	0.280
Adjusted R-squared	0.346	0.828	0.279
AIC	20069.451	20069.451	20377.611
BIC	20136.213	20136.213	20420.096

As we can notice, the value of  $R^2$ , along with the value  $R_{adj}^2$ , differs significantly (compering Model 1 and Model 2), which may lead us to the inaccurate conclusion that Model 2 has stronger explanatory power than Model 1. When we excluded intercept, we faced two practical drawbacks and one of them especially noticeable in this case. Namely, regression packages usually change the way R-squared is computed when an overall intercept is not included by simply replacing SST with a total sum of squares that does not center  $y_i$  about its mean ( $SST_0 = \sum_{i=1}^n y_i^2$ ), so-called uncentered R-squared. Moreover, the resulting R-squared, say  $R_0^2 = 1 - SSR/SST_0$ , is rarely suitable as a goodness of fit measure, as its high value is an artifact of non centering the total sum of squares in the calculation. However, when compering the measures related to “information criteria,” AIC and SC(BIC), their identical value tell us that the out-of-the-sample performance of Model 1 and Model 2 is same, thus we would be indifferent if we chose between these two models.

In line with everything previously said, the higher value of AIC/BIC for Model 3, along with lower value of  $R^2_{adj}$ , makes it to be inferior choice in comparison to Model 1. We would end up with the exact conclusion when comparing Model2 and Model3, as the values of AIC and BIC are the same as for Model1.

We obtain “information criteria” AIC and SC(BIC) using following formula which includes complexity of out-of-sample performance:

$$AIC/BIC = \log\left(\frac{SSR}{n}\right) + \frac{m}{n}(k-1)$$

where for AIC  $m = 2$  and for BIC  $m = \log(n)$ . BIC is especially useful tool/measure if we are building a model with aim of making predictions, as it penalizes increase in  $k$  more than the AIC.

## 2.6

## 2.7

- a)
- b)
- c)
- d)

## 3 Theory

### 3.1

In a general sense,  $SSR$  or Sum of Squared Residuals is a good measure of variation that is *not explained* by the given model. However, as models become more complex, the drawbacks of this measure are getting visible and more serious. The main disadvantage of this measure is that  $SSR$  automatically decreases when the number of explanatory variables increase. Consequently, it affects  $R^2$  in an opposite way bearing in mind the relation between those measures:  $R^2 = 1 - \frac{SSR}{SST}$ . Thus, just relying on the values of the  $SSR$  or  $R^2$ , we would always choose to have as much variables in our model as we can. Unfortunately, by following this way of reasoning we could fall into the trap. Namely, we could face the *bias-variance trade off* which states that higher number of variables can definitely increase the precision of the model inside of the sample, but not necessarily out of the sample.

All things considered, one way to overcome the above mentioned problem is to include additional criteria based on which we can choose the model with *optimal* number of variables. *BIC* criterion is defined as following:

$$BIC = \log\left(\frac{SSR}{N}\right) + \frac{m}{N}(k-1)$$

where  $m = \log(n)$  Therefore, the main purpose of *BIC* is to decide which model can best describe  $Y$  with respect to the number of variables.

In our particular example, considering the possibility of having different explanatory variables in given models, we can not state the same values for  $R^2$  and *BIC* in both models, even though the response variable  $Y$  is the same. However, they will definitely be consistent, meaning that if the conclusion based on  $R^2$  is that one model better explained  $Y$  than another, *BIC* criterion would state the same. The main reason is that models are different just in terms of explanatory variables but with *same amount of them* ( $k_1 = k_2 = k$ ). To rephrase it, in this situation additional criterion will state the same as  $R^2$  because *BIC* is valuable only when we have to go for one from two models, with respect to *different* values of  $k$  or explanatory variables.

To prove our statements mathematically, we should consider following formulas :

$$\text{Model 1: } R_1^2 = 1 - \frac{SSR_1}{SST}, \quad BIC_1 = \log\left(\frac{SSR_1}{N}\right) + \frac{m}{N}(k-1)$$

$$\text{Model 2: } R_2^2 = 1 - \frac{SSR_2}{SST}, \quad BIC_2 = \log\left(\frac{SSR_2}{N}\right) + \frac{m}{N}(k-1)$$

In view of models given in the task, both  $R^2$  and  $BIC$  are identical in terms of all variables in equations except  $SSR$ . Therefore, the model with the smaller value of  $SSR$  or higher value of  $R^2$  will have better in-sample performance. In addition, smaller value of  $SSR$  will have an influence on  $BIC$  as well, lowering the observed value. Having in mind that we want  $BIC$  value to be small in order to choose desired model, both measures lead us to the same, consistent conclusion: model with smaller  $SSR$  will be preferable using both criteria.

With the aim of proving what has been written before, we used the given dataset and run regression of rating on price and education, denoted Model1, and, as Model2, run regression of rating on age and income. The given requirements are met as we are having the same number of observations and explanatory variables.

```
lin.model1 <- lm(rating ~ price + education, data)
lin.model2 <- lm(rating ~ age + income, data)
```

Table 7: Proof for 3.1.

	Model 1	Model 2
R-squared	0.278	0.003
Adjusted R-squared	0.277	0.003
AIC	20382.827	21411.203
BIC	20407.105	21435.481

## 3.2

True: In the regression model

$$Y = \beta_0 + \beta_1 D + \beta_2 D^2 + u, \quad \mathbb{E}[u|D] = 0$$

a quadratic term involves squaring the predictor, where the predictor is dummy-coded variable.

Generally, when it comes to Quadratic regression models, the rule is that perfect multicollinearity assumption is not violated. The reason is because violation of the assumption implies the presence of exact linear relation between any predictors  $X_i$  and  $X_j$  when  $i \neq j$ . Consequently, even though  $X_2 = X_1^2$  while  $X_1 = X$  in example  $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + u$ , variable  $X_2$  can not be written as a linear function of  $X_1$  but rather non-linear, hence multicollinearity assumption still holds. However, in this particular example, when the explanatory variables are *dummy variables*, meaning that they can take either values 1 or 0, the general rule for Quadratic regression models is violated. Namely, having in mind that taking 0 or 1 to the power of 2 does not make any difference in comparison to non-squared dummy variable, this model should be considered as a Multiple regression model rather than the Quadratic. Thus, we will have our perfect multicollinearity assumption to be violated which bring us to the following consequences :

- (1) The matrix  $X'X$  is not invertible, hence the OLS estimator does not exist
- (2) There are infinitely many parameter values  $\gamma$  having the same minimal sum of squared residuals( $SSR(\gamma)$ )
- (3) The parameters in the regression model are not identified (not available: NA)

With the aim of proving what has been written before, we used the given dataset and run regression of rating on education and where education is a dummy variable which takes value of one if consumer has 1 high school diploma or higher.

```

model4 <- lm(rating ~ education + I(education^2), data)

stargazer(model4, header = FALSE, title="Model 4 - Multiple OLS Regression Model")

##
## \begin{table}[!htbp] \centering
##   \caption{Model 4 - Multiple OLS Regression Model}
##   \label{}
## \begin{tabular}{@{\extracolsep{5pt}}lc}
## \hline
## & \multicolumn{1}{c}{\textit{Dependent variable:}} \\
## \cline{2-2}
## & rating \\
## \hline
## education & $-0.493$^{*} \\
## & (0.255) \\
## & \\
## I(education$\hat{\text{ }}$) & \\
## & \\
## & \\
## Constant & 11.872$^{***}$ \\
## & (0.204) \\
## & \\
## \hline
## Observations & 3,195 \\
## R$^{2}$ & 0.001 \\
## Adjusted R$^{2}$ & 0.001 \\
## Residual Std. Error & 6.903 (df = 3193) \\
## F Statistic & 3.745$^{*}$ (df = 1; 3193) \\
## \hline
## \textit{Note:} & \multicolumn{1}{r}{$^{*}\text{p}<\$0.1$; $^{**}\text{p}<\$0.05$; $^{***}\text{p}<\$0.01$} \\
## \end{tabular}
## \end{table}

```

Table 8 was created using stargazer v.5.2.2 by Hlavac (2018).

### 3.3

We are considering the following quadratic model:

$$Y = \beta_0 + \beta_1 X + \beta_2 X^2 + u$$

where  $\mathbb{E}[u|X] = 0$ .

The null hypothesis that states that the point where the effect of a marginal increase in  $X$  on the conditional expectation  $\mathbb{E}[Y|X]$  changes its sign is 1 (in other words, **the value of vertex** equals one) would be written as follows:

$$H_0 : \frac{-\beta_1}{2\beta_2} = 1 \quad \rightarrow \quad H_0 : \beta_1 + 2\beta_2 = 0$$

$$H_1 : \beta_1 + 2\beta_2 \neq 0$$

The distribution of  $\beta_1 + 2\beta_2$  under the null is given by the following univariate normal distribution:

$$\beta_1 + 2\beta_2 | X \sim \mathcal{N}(0, \text{Var}(\beta_1 + 2\beta_2 | X))$$

where  $\text{Var}(\hat{\beta}_1 + 2\hat{\beta}_2) = \text{Var}(\hat{\beta}_1) + 4\text{Var}(\hat{\beta}_2) + 4\text{Cov}(\hat{\beta}_1, \hat{\beta}_2)$ .

The test statistic

$$t = \frac{\hat{\beta}_1 + 2\hat{\beta}_2}{\text{se}(\hat{\beta}_1 + 2\hat{\beta}_2)} | X \sim t_{df}$$

follows the  $t_{df}$  distribution with  $df = (N - K - 1)$  when  $\sigma^2$  is substituted by  $\hat{\sigma}^2$ .

However, the test statistic requires the standard error which is not reported in a standard regression.

The needed standard error is given by the square root of the variance:

$$\text{se}(\hat{\beta}_1 + 2\hat{\beta}_2) = \sqrt{\text{Var}(\hat{\beta}_1 + 2\hat{\beta}_2)}$$

The variance is given by:

$$\text{Var}(\hat{\beta}_1 + 2\hat{\beta}_2) = \text{Var}(\hat{\beta}_1) + 4\text{Var}(\hat{\beta}_2) + 4\text{Cov}(\hat{\beta}_1, \hat{\beta}_2)$$

The estimator of the standard error is thus given by:

$$\text{se}(\hat{\beta}_1 + 2\hat{\beta}_2) = \sqrt{[\text{se}(\hat{\beta}_1)]^2 + [\text{se}(2\hat{\beta}_2)]^2 + 2s_{12}}$$

The sample covariance  $s_{12}$  of the parameters is not given in a standard regression. To get the standard error of a linear combination you either need:

- (1) To rewrite the model so that the needed standard error is given in a standard regression.
- (2) A statistical software that computes the sample covariance between the parameters
- (3) **A statistical software that allows direct testing of linear combinations.**

Under these circumstances, it is crucial to pay attention to the position of the vertex relative to the range of data X, meaning that if the range includes the vertex, the sign would change.

## 3.4

```
set.seed(1)

b0 <- 2
b1 <- -1.5
b2 <- 3
iters <- c(10)

models <- list()
for (i in seq_along(iters)) {
  x <- runif(iters[i], -3, 3)
  x1 <- x^2
```

```

u <- rnorm(iters[i], 0, 2^2)
y <- b0 + b1*x + b2*x1 + u
models[[i]] <- lm(y ~ x + x1) # or: models[[i]] <- lm(y ~ x + I(x^2))
}

stargazer(models[[1]], header = FALSE, title="Simulated Regression Model")

```

Table 8: Simulated Regression Model

<i>Dependent variable:</i>	
	y
x	-1.538* (0.745)
x1	3.872*** (0.497)
Constant	-0.541 (2.073)
Observations	10
R <sup>2</sup>	0.897
Adjusted R <sup>2</sup>	0.867
Residual Std. Error	4.113 (df = 7)
F Statistic	30.350*** (df = 2; 7)

Note: \*p<0.1; \*\*p<0.05; \*\*\*p<0.01

Table 8 was created using stargazer v.5.2.2 by Hlavac (2018).

```
kable(t(range(x)), col.names = c("Min", "Max"), caption = "The range of explanatory variable")
```

Table 9: The range of explanatory variable

Min	Max
-2.629282	2.668052

Note that the vertex is in the range of explanatory variable.

```

kable(linearHypothesis(models[[1]], c("x=-2*x1")),
      caption = "Testing the null hypothesis:
The vertex of a quadratic equation equals one.")

```

Table 10: Testing the null hypothesis: The vertex of a quadratic equation equals one.

Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
8	663.4941	NA	NA	NA	NA
7	118.4034	1	545.0907	32.22572	0.0007532

The answer is p-value = 0.0007532, which fails to reject the null hypothesis. Thus, the conjecture that a point where the effect of a marginal increase in X on the conditional expectation  $\mathbb{E}[Y|X]$  changes its sign is 1 is not supported by the data.

## References

Hlavac, Marek. 2018. *Stargazer: Well-Formatted Regression and Summary Statistics Tables*. Bratislava, Slovakia: Central European Labour Studies Institute (CELSI). <https://CRAN.R-project.org/package=stargazer>.